

Midterm Review

Hongning Wang

CS@UVa

Core concepts

- Search Engine Architecture
 - Key components in a modern search engine
- Crawling & Text processing
 - Different strategies for crawling
 - Challenges in crawling
 - Text processing pipeline
 - Zipf's law
- Inverted index
 - Index compression
 - Phrase queries

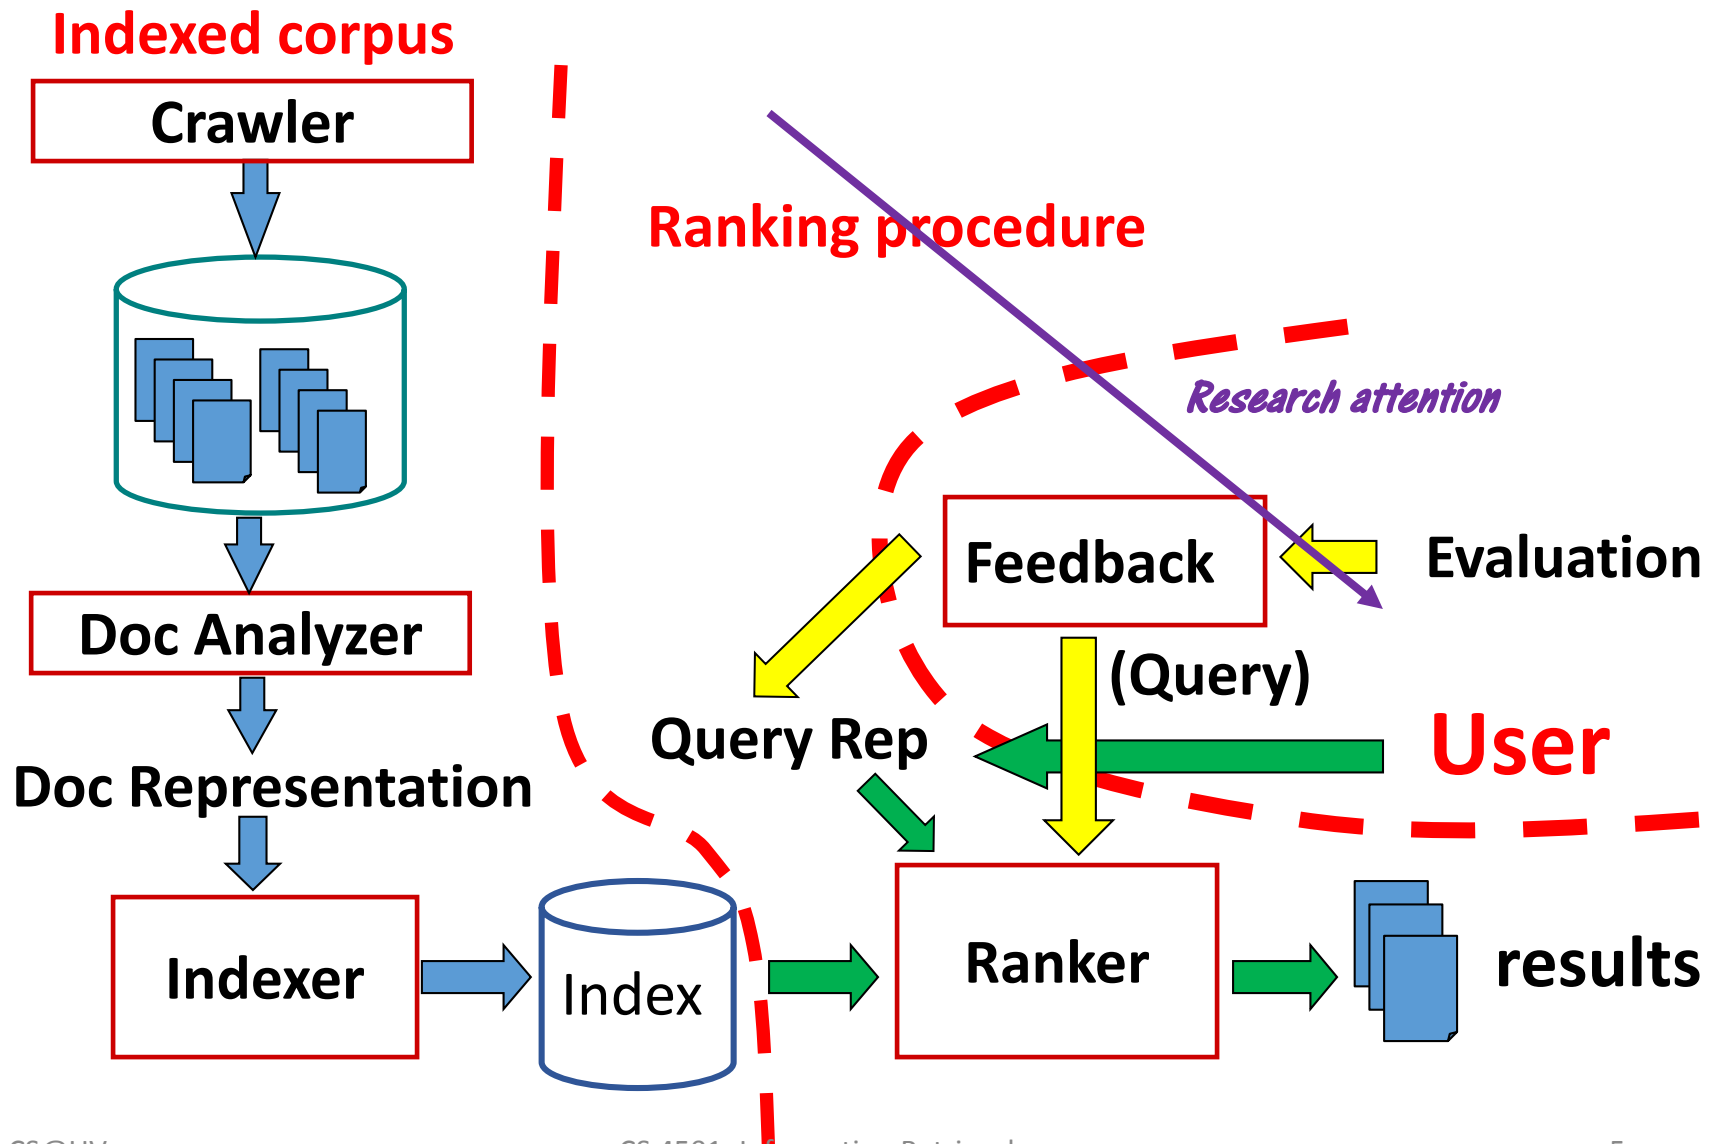
Core concepts

- Vector space model
 - Basic term/document weighting schemata
- Latent semantic analysis
 - Word space to concept space
- Probabilistic ranking principle
 - Risk minimization
 - Document generation model
- Language model
 - N-gram language models
 - Smoothing techniques

Core concepts

- Classical IR evaluations
 - Basic components in IR evaluations
 - Evaluation metrics
 - Annotation strategy and annotation consistency

Abstraction of search engine architecture



IR v.s. DBs

- Information Retrieval:
 - Unstructured data
 - Semantics of objects are subjective
 - Simple keyword queries
 - Relevance-drive retrieval
 - Effectiveness is primary issue, though efficiency is also important
- Database Systems:
 - Structured data
 - Semantics of each object are well defined
 - Structured query languages (e.g., SQL)
 - Exact retrieval
 - Emphasis on efficiency

Crawler: visiting strategy

- Breadth first
 - Uniformly explore from the entry page
 - Memorize all nodes on the previous level
 - As shown in pseudo code
- Depth first
 - Explore the web by branch
 - Biased crawling given the web is not a tree structure
- Focused crawling
 - Prioritize the new links by predefined strategies
- Challenges
 - Avoid duplicate visits
 - Re-visit policy

Automatic text indexing

- Tokenization
 - Regular expression based
 - Learning-based
- Normalization
- Stemming
- Stopword removal

Statistical property of language


discrete version of power law

- Zipf's law
 - Frequency of any word is inversely proportional to its rank in the frequency table
 - Formally
 - $f(k; s, N) = \frac{1/k^s}{\sum_{n=1}^N 1/n^s}$
where k is rank of the word; N is the vocabulary size; s is language-specific parameter

Automatic text indexing

Remove non-informative words

Remove 1s

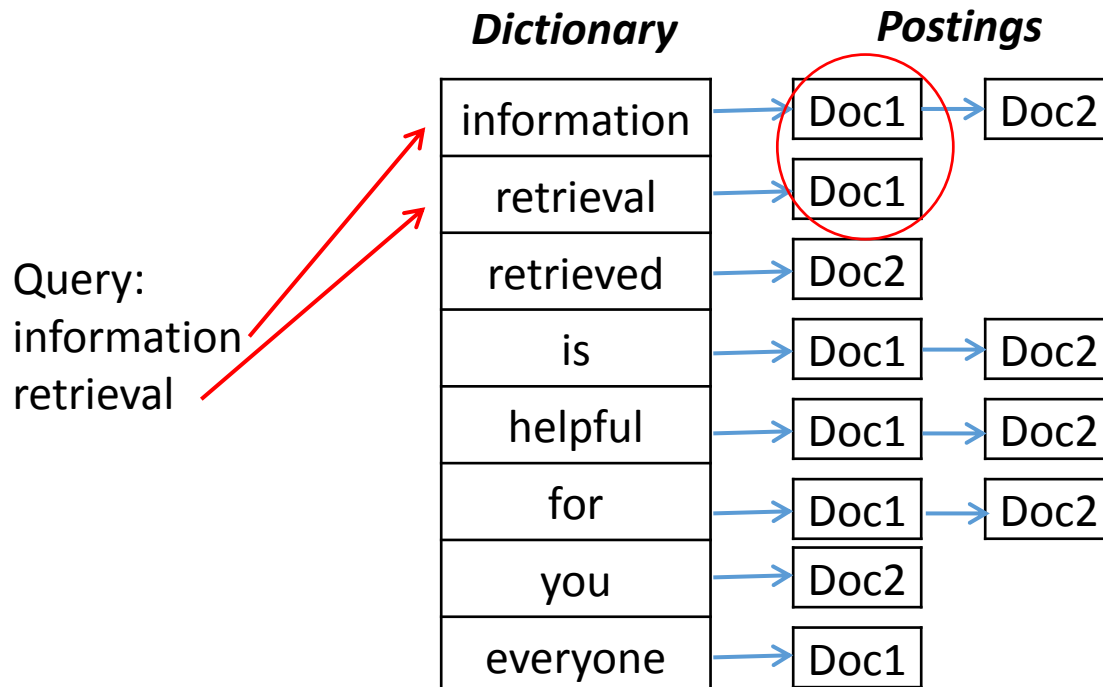


Terms	Documents													
	M1	M2	M3	M4	M5	M6	M7	M8	M9	M10	M11	M12	M13	M14
abnormalities	0	0	0	0	0	0	0	1	0	1	0	0	0	0
age	1	0	0	0	0	0	0	0	0	0	0	1	0	0
behavior	0	0	0	0	1	1	0	0	0	0	0	0	0	0
blood	0	0	0	0	0	0	0	1	0	0	1	0	0	0
close	0	0	0	0	0	0	1	0	0	0	1	0	0	0
culture	1	1	0	0	0	0	0	1	1	0	0	0	0	0
depressed	1	0	1	1	1	0	0	0	0	0	0	0	0	0
discharge	1	1	0	0	0	1	0	0	0	0	0	0	0	0
disease	0	0	0	0	0	0	0	0	1	0	1	0	0	0
fast	0	0	0	0	0	0	0	0	0	1	0	1	1	1
generation	0	0	0	0	0	0	0	0	1	0	0	0	1	0
oestrogen	0	0	1	1	0	0	0	0	0	0	0	0	0	0
patients	1	1	0	1	0	0	0	1	0	0	0	0	0	0
pressure	0	0	0	0	0	0	0	0	0	0	1	0	0	1
rats	0	0	0	0	0	0	0	0	0	0	0	0	1	1
respect	0	0	0	0	0	0	0	1	0	0	0	1	0	0
rise	0	0	0	1	0	0	0	0	0	0	0	0	0	1
study	1	0	1	0	0	0	0	0	1	0	0	0	0	0

Figure 2.1. A plot of the hyperbolic curve relating f , the frequency of occurrence and r , the rank order (Adapted from Schultz⁴⁴ page 127)

Inverted index

- Build a look-up table for each word in vocabulary
 - From word to find documents!



Time complexity:

- $O(|q| * |L|)$, $|L|$ is the average length of posting list
- By Zipf's law, $|L| \ll D$

Structures for inverted index

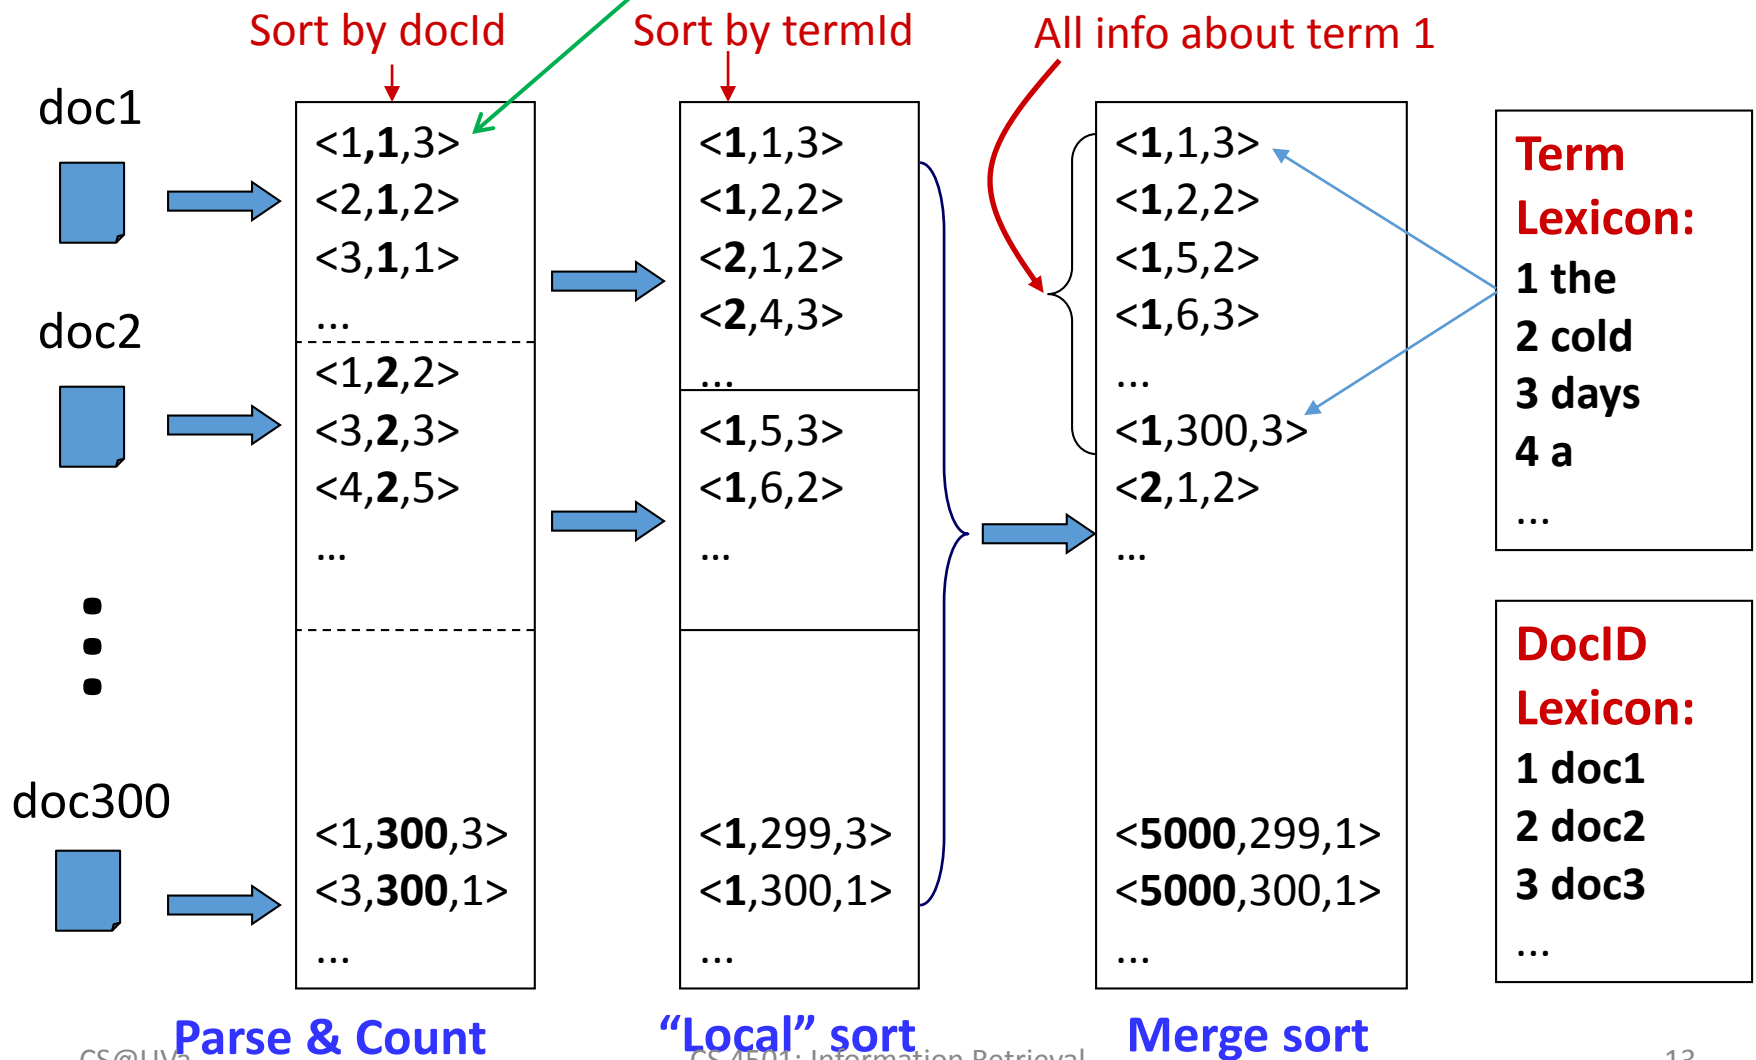
- Dictionary: modest size
 - Needs fast random access
 - Stay in memory
 - Hash table, B-tree, trie, ...
- Postings: huge
 - Sequential access is expected
 - Stay on disk
 - Contain docID, term freq, term position, ...
 - Compression is needed

“Key data structure underlying modern IR”

- Christopher D. Manning

Sorting-based inverted index construction

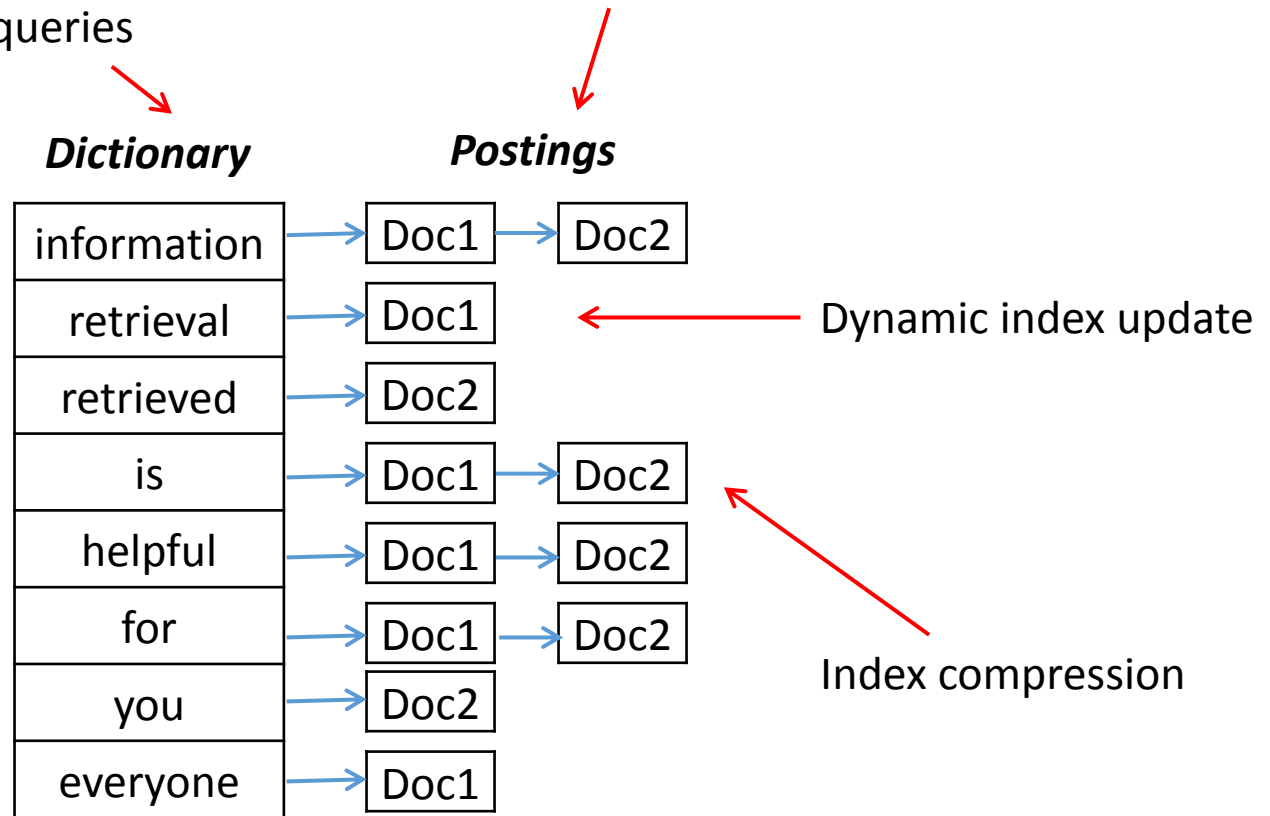
<Tuple>: <termID, docID, count>



A close look at inverted index

Approximate search:
e.g., misspelled queries,
wildcard queries

Proximity search:
e.g., phrase queries




Index compression

- Observation of posting files
 - Instead of storing docID in posting, we store gap between docIDs, since they are ordered
 - Zipf's law again:
 - The more frequent a word is, the smaller the gaps are
 - The less frequent a word is, the shorter the posting list is
 - Heavily biased distribution gives us great opportunity of compression!

Information theory: entropy measures compression difficulty.

Phrase query

- Generalized postings matching
 - Equality condition check with requirement of position pattern between two query terms
 - e.g., $T2.pos - T1.pos = 1$ (T1 must be immediately before T2 in any matched document)
 - Proximity query: $|T2.pos - T1.pos| \leq k$



Considerations in result display

- Relevance
 - Order the results by relevance
- Diversity
 - Maximize the topical coverage of the displayed results
- Navigation
 - Help users easily explore the related search space
 - Query suggestion
 - Search by example

Deficiency of Boolean model

- The query is unlikely precise
 - “Over-constrained” query (terms are too specific): no relevant documents can be found
 - “Under-constrained” query (terms are too general): over delivery
 - It is hard to find the right position between these two extremes (hard for users to specify constraints)
- Even if it is accurate
 - Not all users would like to use such queries
 - All relevant documents are **not equally** important
 - No one would go through all the matched results
- Relevance is a matter of degree!

Vector space model

- Represent both doc and query by concept vectors
 - Each concept defines one dimension
 - K concepts define a high-dimensional space
 - Element of vector corresponds to concept weight
 - E.g., $d=(x_1, \dots, x_k)$, x_i is “importance” of concept i
- Measure relevance
 - Distance between the query vector and document vector in this concept space


What is a good “basic concept”?

- Orthogonal
 - Linearly independent basis vectors
 - “Non-overlapping” in meaning
 - No ambiguity
- Weights can be assigned automatically and accurately
- Existing solutions
 - Terms or N-grams, i.e., bag-of-words
 - Topics, i.e., topic model

TF normalization

- Sublinear TF scaling

- $tf(t, d) = \begin{cases} 1 + \log f(t, d), & \text{if } f(t, d) > 0 \\ 0, & \text{otherwise} \end{cases}$




TF normalization

- Maximum TF scaling

- $tf(t, d) = \alpha + (1 - \alpha) \frac{f(t, d)}{\max_t f(t, d)}$

- Normalize by the most frequent word in this doc



IDF weighting

- Solution

- Assign higher weights to the rare terms

- Formula

- $IDF(t) = 1 + \log\left(\frac{N}{df(t)}\right)$

Non-linear scaling

Total number of docs in collection

Number of docs containing term t

- A corpus-specific property

- Independent of a single document

TF-IDF weighting

- Combining TF and IDF
 - Common in doc \rightarrow high tf \rightarrow high weight
 - Rare in collection \rightarrow high idf \rightarrow high weight
 - $w(t, d) = TF(t, d) \times IDF(t)$
- Most well-known document representation schema in IR! (G Salton et al. 1983)




"Salton was perhaps the leading computer scientist working in the field of information retrieval during his time." - wikipedia

[Gerard Salton Award](#)

– highest achievement award in IR

From distance to angle

- Angle: how vectors are overlapped
 - Cosine similarity – projection of one vector onto another



Advantages of VS Model

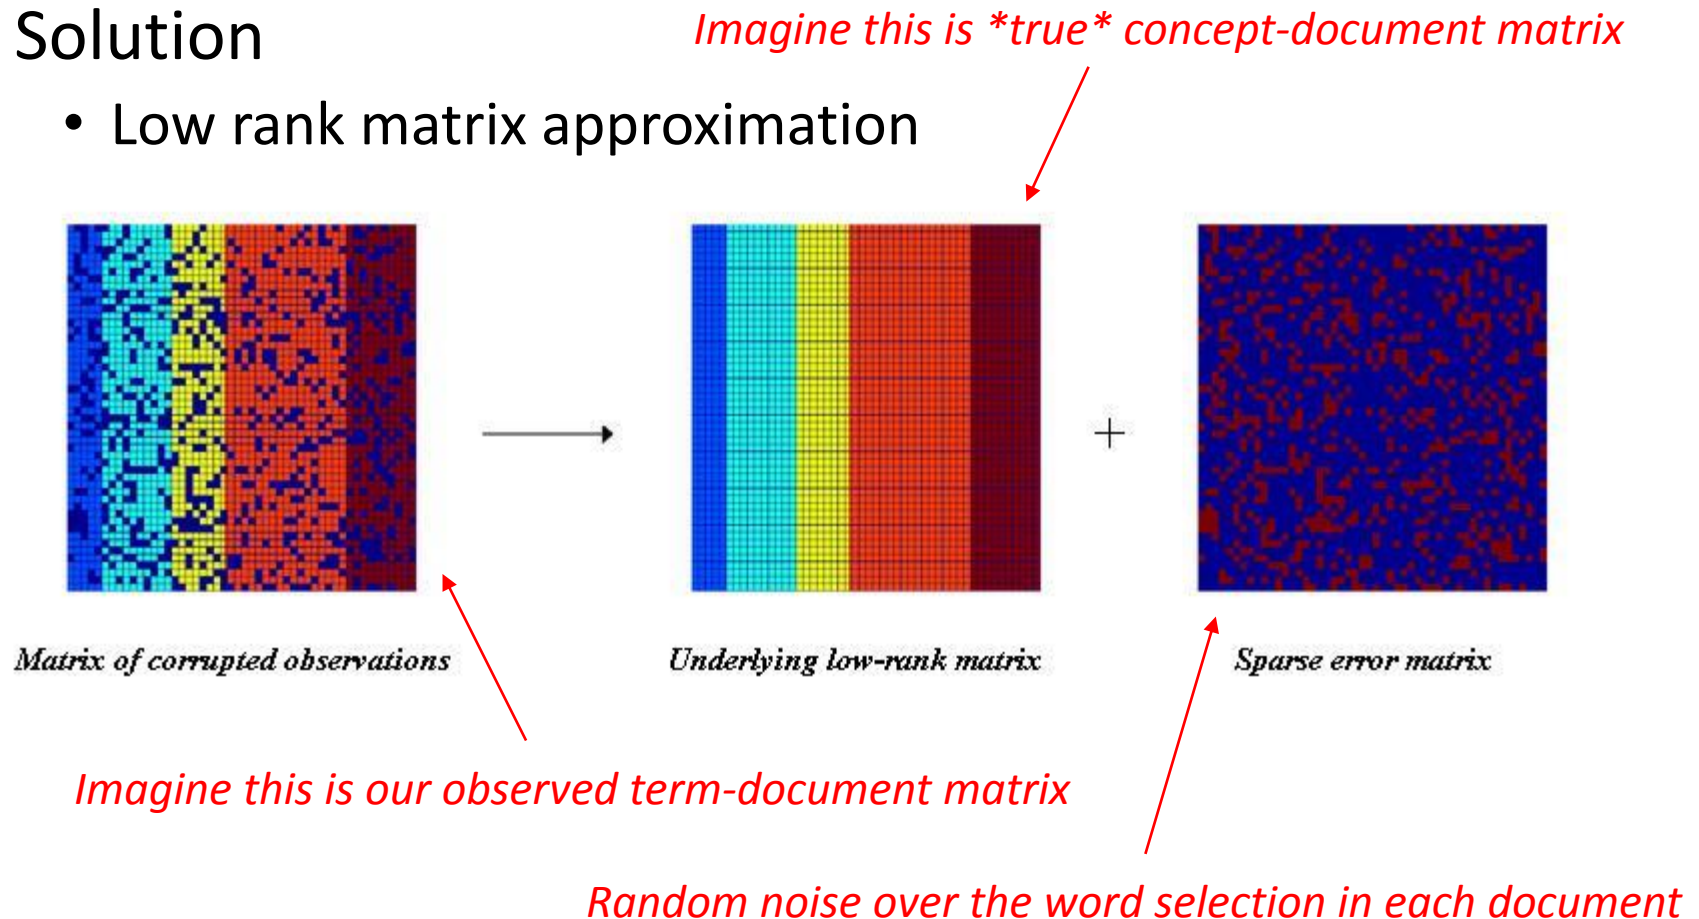
- Empirically effective! (Top TREC performance)
- Intuitive
- Easy to implement
- Well-studied/Mostly evaluated
- The Smart system
 - Developed at Cornell: 1960-1999
 - Still widely used
- **Warning: Many variants of TF-IDF!**

Disadvantages of VS Model

- Assume term independence
- Assume query and document to be the same
- Lack of “predictive adequacy”
 - Arbitrary term weighting
 - Arbitrary similarity measure
- Lots of parameter tuning!


Latent semantic analysis

- Solution
 - Low rank matrix approximation



Latent Semantic Analysis (LSA)

- Solve LSA by SVD



1. Perform SVD on document-term adjacency matrix
2. Construct $C_{M \times N}^k$ by only keeping the largest k singular values in Σ non-zero

Probabilistic ranking principle

- From decision theory
 - Two types of loss
 - $\text{Loss}(\text{retrieved} | \text{non-relevant}) = a_1$
 - $\text{Loss}(\text{not retrieved} | \text{relevant}) = a_2$
 - $\phi(d_i, q)$: probability of d_i being relevant to q
 - Expected loss regarding to the decision of including d_i in the final results
 - Retrieve: $(1 - \phi(d_i, q))a_1$
 - Not retrieve: $\phi(d_i, q)a_2$

Your decision criterion?

Probabilistic ranking principle

- From decision theory
 - We make decision by
 - If $(1 - \phi(d_i, q))a_1 < \phi(d_i, q)a_2$, retrieve d_i
 - Otherwise, not retrieve d_i
 - Check if $\phi(d_i, q) > \frac{a_1}{a_1 + a_2}$
 - Rank documents by descending order of $\phi(d_i, q)$ would minimize the loss

Conditional models for $P(R=1 | Q, D)$

- Basic idea: relevance depends on how well a query matches a document
 - $P(R=1 | Q, D) = g(\text{Rep}(Q, D) | \theta)$ ← a functional form
 - $\text{Rep}(Q, D)$: feature representation of query-doc pair
 - E.g., #matched terms, highest IDF of a matched term, docLen
 - Using training data (with known relevance judgments) to estimate parameter θ
 - Apply the model to rank new documents
- Special case: logistic regression

Generative models for $P(R=1 | Q, D)$

- Basic idea

- Compute $\text{Odd}(R=1 | Q, D)$ using Bayes' rule

$$\text{Odd}(R=1 | Q, D) = \frac{P(R=1 | Q, D)}{P(R=0 | Q, D)} = \frac{P(Q, D | R=1)}{P(Q, D | R=0)} \frac{P(R=1)}{P(R=0)} \leftarrow \text{Ignored for ranking}$$

- Assumption

- Relevance is a binary variable

- Variants

- Document “generation”
 - $P(Q, D | R) = P(D | Q, R)P(Q | R)$
 - Query “generation”
 - $P(Q, D | R) = P(Q | D, R)P(D | R)$

Document generation model

$$\begin{aligned}
 \text{Odd}(R=1|Q,D) &\propto \prod_{i=1}^k \frac{P(A_i = d_i | Q, R=1)}{P(A_i = d_i | Q, R=0)} \\
 &= \prod_{i=1, d_i=1}^k \frac{P(A_i = 1 | Q, R=1)}{P(A_i = 1 | Q, R=0)} \prod_{i=1, d_i=0}^k \frac{P(A_i = 0 | Q, R=1)}{P(A_i = 0 | Q, R=0)} \\
 &\approx \prod_{i=1, d_i=q_i=1}^k \frac{p_i}{u_i} \prod_{i=1, d_i=0, q_i=1}^k \frac{1-p_i}{1-u_i} \\
 &= \prod_{i=1, d_i=q_i=1}^k \frac{p_i(1-u_i)}{u_i(1-p_i)} \prod_{i=1, q_i=1}^k \frac{1-p_i}{1-u_i}
 \end{aligned}$$

Terms *occur* in doc
 Terms do *not occur* in doc
 Assumption: terms not occurring in the query are equally likely to occur in relevant and nonrelevant documents, i.e., $p_t = u_t$

Important tricks

document	relevant(R=1)	nonrelevant(R=0)
term present $A_i=1$	p_i	u_i
term absent $A_i=0$	$1-p_i$	$1-u_i$

Maximum likelihood estimation

- Data: a document d with counts $c(w_1), \dots, c(w_N)$
- Model: multinomial distribution $p(W|\theta)$ with parameters $\theta_i = p(w_i)$
- Maximum likelihood estimator: $\hat{\theta} = \operatorname{argmax}_{\theta} p(W|\theta)$

$$p(W|\theta) = \binom{N}{c(w_1), \dots, c(w_N)} \prod_{i=1}^N \theta_i^{c(w_i)} \propto \prod_{i=1}^N \theta_i^{c(w_i)} \quad \Rightarrow \quad \log p(W|\theta) = \sum_{i=1}^N c(w_i) \log \theta_i$$

$$\Rightarrow L(W, \theta) = \sum_{i=1}^N c(w_i) \log \theta_i + \lambda \left(\sum_{i=1}^N \theta_i - 1 \right)$$

Using Lagrange multiplier approach, we'll tune θ_i to maximize $L(W, \theta)$

$$\Rightarrow \frac{\partial L}{\partial \theta_i} = \frac{c(w_i)}{\theta_i} + \lambda \rightarrow \theta_i = -\frac{c(w_i)}{\lambda}$$

Set partial derivatives to zero

$$\Rightarrow \text{Since } \sum_{i=1}^N \theta_i = 1 \text{ we have } \lambda = -\sum_{i=1}^N c(w_i)$$

Requirement from probability

$$\Rightarrow \theta_i = \frac{c(w_i)}{\sum_{i=1}^N c(w_i)}$$

The BM25 formula

- A closer look

$$rel(q, D) = \sum_{i=1}^n ID F(q_i) \frac{tf_i(k_1 + 1)}{tf_i + k_1(1 - b + b \frac{|D|}{avg |D|})} \frac{qtf_i(k_2 + 1)}{k_2 + qtf_i}$$


TF-IDF component for document (green arrow pointing to the first fraction)

TF component for query (blue arrow pointing to the second fraction)

- b is usually set to $[0.75, 1.2]$
- k_1 is usually set to $[1.2, 2.0]$
- k_2 is usually set to $(0, 1000]$

**Vector space model
with TF-IDF schema!**

Source-Channel framework [Shannon 48]



$$\hat{X} = \arg \max_X p(X | Y) = \arg \max_X p(Y | X) p(X) \quad (\text{Bayes Rule})$$

When X is text, $p(X)$ is a language model

Many Examples:

Speech recognition:	X =Word sequence	Y =Speech signal
Machine translation:	X =English sentence	Y =Chinese sentence
OCR Error Correction:	X =Correct word	Y = Erroneous word
Information Retrieval:	X =Document	Y =Query
Summarization:	X =Summary	Y =Document

More sophisticated LMs

- N-gram language models
 - In general, $p(w_1 w_2 \dots w_n) = p(w_1)p(w_2|w_1) \dots p(w_n|w_1 \dots w_{n-1})$
 - N-gram: conditioned only on the past N-1 words
 - E.g., bigram: $p(w_1 \dots w_n) = p(w_1)p(w_2|w_1) p(w_3|w_2) \dots p(w_n|w_{n-1})$
- Remote-dependence language models (e.g., Maximum Entropy model)
- Structured language models (e.g., probabilistic context-free grammar)

Justification from PRP

$$\begin{aligned} O(R=1|Q,D) &\propto \frac{P(Q,D|R=1)}{P(Q,D|R=0)} \\ &= \frac{P(Q|D,R=1)P(D|R=1)}{P(Q|D,R=0)P(D|R=0)} \\ &\propto \frac{P(Q|D,R=1)}{P(Q|D,R=0)} \frac{P(D|R=1)}{P(D|R=0)} \quad (\text{Assume } P(Q|D,R=0) \approx P(Q|R=0)) \end{aligned}$$

Query generation (points to $P(D|R=1)$)

Query likelihood $p(q|\theta_d)$ (points to $P(Q|D,R=1)$)

Document prior (points to $P(D|R=1)$)

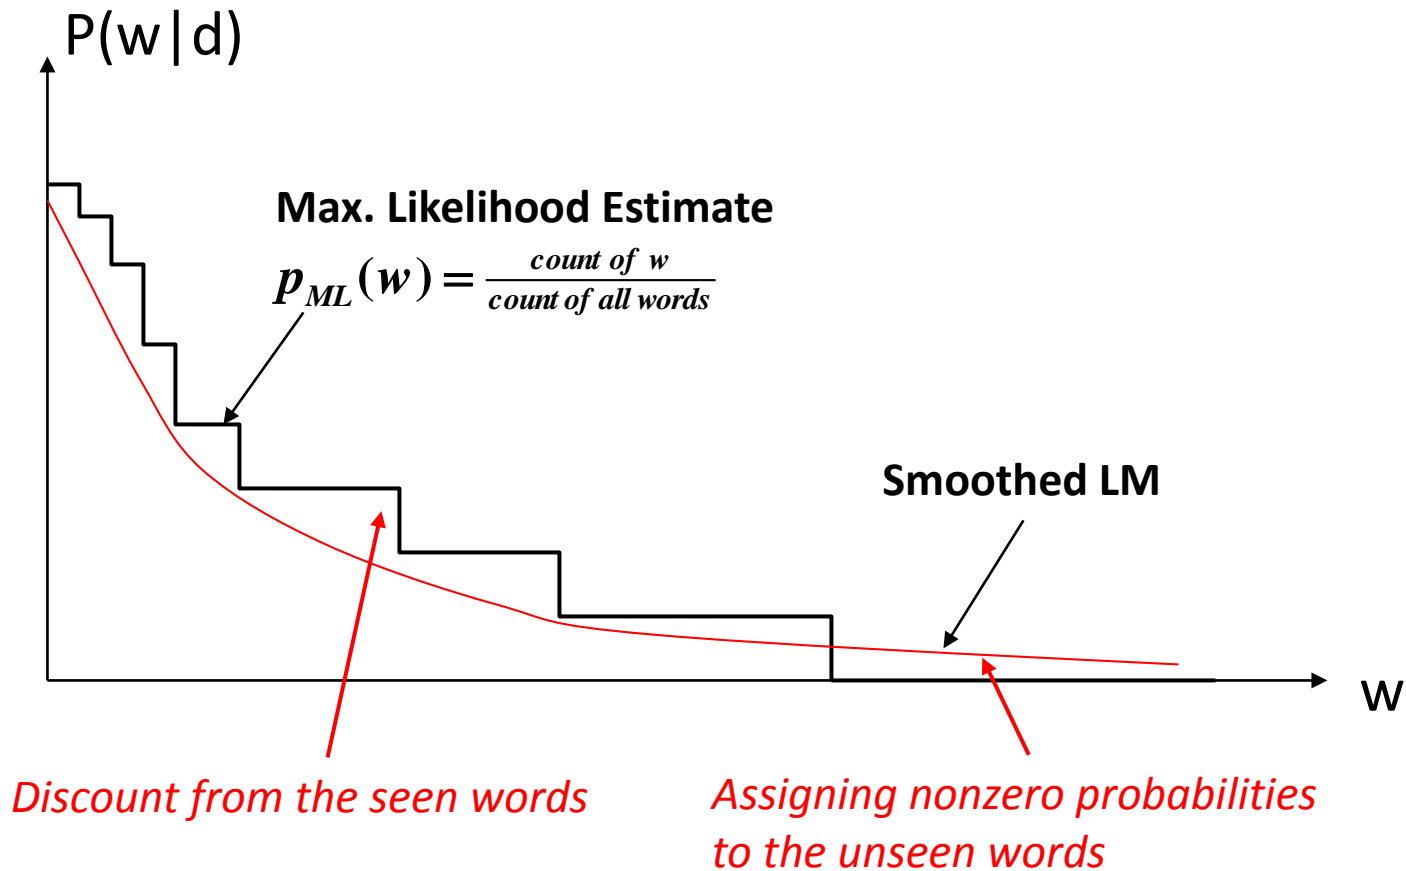
Assuming uniform document prior, we have

$$O(R=1|Q,D) \propto P(Q|D,R=1)$$

Problem with MLE

- What probability should we give a word that has not been observed in the document?
 - $\log 0$?
- If we want to assign non-zero probabilities to such words, we'll have to discount the probabilities of observed words
- This is so-called “smoothing”

Illustration of language model smoothing



Refine the idea of smoothing

- Should all unseen words get equal probabilities?
- We can use a reference model to discriminate unseen words

$$p(w | d) = \begin{cases} p_{seen}(w | d) & \text{if } w \text{ is seen in } d \\ \alpha_d p(w | REF) & \text{otherwise} \end{cases}$$

Discounted ML estimate

Reference language model

$$\alpha_d = \frac{1 - \sum_{w \text{ is seen}} p_{seen}(w | d)}{\sum_{w \text{ is unseen}} p(w | REF)}$$

Smoothing methods

- Method 1: Additive smoothing
 - Add a constant δ to the counts of each word

Counts of w in d

$$p(w | d) = \frac{c(w, d) + 1}{|d| + |V|}$$

“Add one”, Laplace smoothing

Vocabulary size

The diagram shows the formula for Laplace smoothing, $p(w | d) = \frac{c(w, d) + 1}{|d| + |V|}$, highlighted in a yellow box. Three annotations with arrows point to parts of the formula: 'Counts of w in d' points to the numerator term $c(w, d)$; '“Add one”, Laplace smoothing' points to the '+1' in the numerator; and 'Vocabulary size' points to the $|V|$ term in the denominator. An additional arrow points from the text 'Length of d (total counts)' in the following list item to the $|d|$ term in the denominator.

- Problems? **Length of d (total counts)**
 - Hint: all words are equally important?

Smoothing methods

- Method 2: Absolute discounting
 - Subtract a constant δ from the counts of each word

$$p(w|d) = \frac{\max(c(w;d) - \delta, 0) + \delta |d|_u p(w|REF)}{|d|}$$

uniq words

- Problems?
 - Hint: varied document length?

Smoothing methods

- Method 3: Linear interpolation, Jelinek-Mercer
 - “Shrink” uniformly toward $p(w | REF)$

$$p(w | d) = (1 - \lambda) \frac{c(w, d)}{|d|} + \lambda p(w | REF)$$

- Problems?

- Hint: what is missing?

MLE

parameter

Smoothing methods

- Method 4: Dirichlet Prior/Bayesian
 - Assume pseudo counts $\mu p(w | REF)$

$$p(w | d) = \frac{c(w; d) + \mu p(w | REF)}{|d| + \mu} = \frac{|d|}{|d| + \mu} \frac{c(w, d)}{|d|} + \frac{\mu}{|d| + \mu} p(w | REF)$$

- Problems?

parameter


Two-stage smoothing [Zhai & Lafferty 02]

Stage-1

- Explain unseen words
- Dirichlet prior (Bayesian)

Stage-2

- Explain noise in query
- 2-component mixture



$$P(w|d) = (1-\lambda) \frac{c(w,d) + \underbrace{\mu p(w|C)}_{\text{Collection LM}}}{|d| + \underbrace{\mu}_{\text{User background model}}} + \lambda p(w|U)$$

Collection LM

User background model

Understanding smoothing

Topical words

	Query = “the	algorithms	for	data	mining”	
$p_{ML}(w d1):$	0.04	0.001	0.02	0.002	0.003	4.8×10^{-12}
$p_{ML}(w d2):$	0.02	0.001	0.01	0.003	0.004	2.4×10^{-12}

$p(\text{“algorithms”} | d1) = p(\text{“algorithms”} | d2)$
 $p(\text{“data”} | d1) < p(\text{“data”} | d2)$
 $p(\text{“mining”} | d1) < p(\text{“mining”} | d2)$

Intuitively, d2 should have a higher score, but $p(q|d1) > p(q|d2)$...

So we should make $p(\text{“the”})$ and $p(\text{“for”})$ **less different** for all docs, and smoothing helps to achieve this goal...

After smoothing with $p(w|d) = 0.1p_{DML}(w|d) + 0.9p(w|REF)$, $p(q|d1) < p(q|d2)$!

Query	= “the	algorithms	for	data	mining”
$P(w REF)$	0.2	0.00001	0.2	0.00001	0.00001
Smoothed $p(w d1):$	0.184	0.000109	0.182	0.000209	0.000309
Smoothed $p(w d2):$	0.182	0.000109	0.181	0.000309	0.000409

Smoothing & TF-IDF weighting

Smoothed ML estimate

Retrieval formula using the general smoothing scheme

$$p(w|d) = \begin{cases} p_{Seen}(w|d) & \text{if } w \text{ is seen in } d \\ \alpha_d p(w|C) & \text{otherwise} \end{cases}$$



$$\alpha_d = \frac{1 - \sum_{w \text{ is seen}} p_{Seen}(w|d)}{\sum_{w \text{ is unseen}} p(w|C)}$$

Reference language model

$$\begin{aligned} \log p(q|d) &= \sum_{w \in V, c(w,q) > 0} c(w,q) \log p(w|d) \\ &= \sum_{\substack{w \in V, c(w,d) > 0, \\ c(w,q) > 0}} c(w,q) \log p_{Seen}(w|d) + \sum_{w \in V, c(w,q) > 0, c(w,d) = 0} c(w,q) \log \alpha_d p(w|C) \\ &= \sum_{\substack{w \in V, c(w,d) > 0, \\ c(w,q) > 0}} c(w,q) \log p_{Seen}(w|d) + \sum_{w \in V, c(w,q) > 0} c(w,q) \log \alpha_d p(w|C) - \sum_{w \in V, c(w,q) > 0, c(w,d) > 0} c(w,q) \log \alpha_d p(w|C) \\ &= \sum_{\substack{w \in V, c(w,d) > 0, \\ c(w,q) > 0}} c(w,q) \log \frac{p_{Seen}(w|d)}{\alpha_d p(w|C)} + |q| \log \alpha_d + \sum_{w \in V, c(w,q) > 0} c(w,q) \log p(w|C) \end{aligned}$$

Key rewriting step (where did we see it before?)

Similar rewritings are very common when using probabilistic models for IR...

Retrieval evaluation

- Aforementioned evaluation criteria are all good, but not essential
 - Goal of any IR system
 - Satisfying users' information need
 - Core quality measure criterion
 - *“how well a system meets the information needs of its users.” – wiki*
 - Unfortunately vague and hard to execute

Classical IR evaluation

- Three key elements for IR evaluation
 1. A document collection
 2. A test suite of information needs, expressible as queries
 3. A set of relevance judgments, e.g., binary assessment of either *relevant* or *nonrelevant* for each query-document pair

Evaluation of unranked retrieval sets

- Summarizing precision and recall to a single value
 - In order to compare different systems
 - F-measure: weighted harmonic mean of precision and recall, α balances the trade-off


$$F = \frac{1}{\alpha \frac{1}{P} + (1 - \alpha) \frac{1}{R}} \quad \left(F_1 = \frac{2}{\frac{1}{P} + \frac{1}{R}} \right)$$

- Why harmonic mean?

- System1: P:0.53, R:0.36
- System2: P:0.01, R:0.99

H	A
0.429	0.445
0.019	0.500

Equal weight between precision and recall



Evaluation of ranked retrieval results

- Summarize the ranking performance with a single number
 - Binary relevance
 - Eleven-point interpolated average precision
 - Precision@K (P@K)
 - Mean Average Precision (MAP)
 - Mean Reciprocal Rank (MRR)
 - Multiple grades of relevance
 - Normalized Discounted Cumulative Gain (NDCG)

AvgPrec is about one query

 = the relevant documents





Figure from Manning Stanford CS276, Lecture 8

AvgPrec of the two rankings

$$\text{Ranking \#1: } (1.0 + 0.67 + 0.75 + 0.8 + 0.83 + 0.6) / 6 = 0.78$$

$$\text{Ranking \#2: } (0.5 + 0.4 + 0.5 + 0.57 + 0.56 + 0.6) / 6 = 0.52$$

What does query averaging hide?




Figure from Doug Oard's presentation, originally from Ellen Voorhees' presentation

Measuring assessor consistency

- *kappa* statistic

- A measure of agreement between judges

$$\kappa = \frac{P(A) - P(E)}{1 - P(E)}$$

- $P(A)$ is the proportion of the times judges agreed
- $P(E)$ is the proportion of times they would be expected to agree by chance
- $\kappa = 1$ if two judges always agree
- $\kappa = 0$ if two judges agree by chance
- $\kappa < 0$ if two judges always disagree

What we have not considered

- The physical form of the output
 - User interface
- The effort, intellectual or physical, demanded of the user
 - User effort when using the system
- Bias IR research towards optimizing relevance-centric metrics